

ESA Gap And Novelty Audit

The attached ESA PRD defines a credible thesis, but without repository or runtime artifacts in the provided evidence, the current implementation cannot honestly be scored as complete; most implementation-dependent claims are **UNVERIFIED** and several differentiators remain **DOCUMENTED ONLY** rather than proven in code or demo behavior (). The attached PRD itself is unusually clear about the intended authority split—"AI should propose bounded infrastructure actions... deterministic controls should verify, modify, deny, execute, and audit those actions" and "A deterministic safety layer remains the final authority" (). If Razorpay engineers saw a demo that visibly enforced that split through typed contracts, stale-state rejection, a single Action Gateway, real topology mutation, and replayable audit lineage, they would likely find it worth exploring; if they instead saw agents recommending scale-out into a policy service that mostly always allows and a UI that simulates recovery, they would classify it as standard agentic AIOps wrapped in payment language (Fatmi, 2026; Vadapalli, 2026).

Evidence Coverage

Evidence Coverage Across ESA

	Payments	Kubernetes	MAPE-K	Agent Governance	Intent Loops
Authority Boundary	GAP	GAP	GAP	GAP	GAP
State Versioning	GAP	GAP	GAP	GAP	GAP
Effect Verification	GAP	GAP	GAP	GAP	GAP
Replayability	GAP	GAP	GAP	GAP	GAP
Benchmarking	GAP	GAP	GAP	GAP	GAP

	Payments	Kubernetes	MAPE-K	Agent Governance	Intent Loops
The strongest coverage in the supplied literature is around closed-loop adaptation, policy-bounded execution, and agent governance, not around ESA-specific “executable state” as a mature implemented pattern (Arcaini et al., 2015; Gomes et al., 2021; Fatmi, 2026; Sohail & Haider, 2026). The biggest gap is the combination ESA wants to claim: one system that joins payment-aware workload state, version-aware action validity, deterministic authority, post-action effect measurement, and replayable decisions in a public end-to-end prototype; adjacent papers usually cover subsets rather than the whole package (Fatmi, 2026; He & Yu, 2026; Ganapathi & Gupta, 2026).					

FIGURE 1 Research coverage across ESA architectural dimensions

Executive Verdict

ESA is not yet defensibly “complete” from the evidence shown here, because the attached source is a PRD, not code, and its own success criteria are framed as future implementation checkpoints rather than observed achieved results (). The PRD explicitly targets one complete measurable loop, at least two typed actions, one blocked action, one stale-state rejection, visible topology change, rollback, and a rule-only benchmark; until runtime proof exists for those, the safest statuses are **DOCUMENTED ONLY** for architecture and **UNVERIFIED** for behavior ().

The interesting part is narrower than the PRD’s broader language. Multi-agent monitoring, diagnosis, planning, and remediation are already well represented in cloud and Kubernetes literature, including reproducible synthetic-workload evaluations, observability-driven healing, typed tool layers, and automated mitigation loops (Avgerinos et al., 2026; Aravindakshan, 2026; Sannareddy, 2025; Vadapalli, 2026). ESA becomes interesting only where it pushes beyond that baseline into a non-bypassable execution boundary with state/version-aware validation and replayable decision artifacts (Fatmi, 2026; He & Yu, 2026; Rhodes & Kang, 2026; Fernandez, 2026).

PRD Compliance

Requirement	Status	Evidence	Exact Gap	Priority
Event ingestion + synthetic workload	DOCUMENTED ONLY	Attached PRD requires synthetic generator and optional Razorpay Test Mode/webhooks ()	No repository/runtime proof of actual ingest path	P1
Workload state/state fabric	DOCUMENTED ONLY	PRD defines WorkloadState, shard state, topology, policies, and metadata ()	No evidence state fabric persists or drives execution	P1

Requirement	Status	Evidence	Exact Gap	Priority
State versioning	DOCUMENTED ONLY	“Every transition is versioned”; stale actions must be rejected ()	No proof version checked atomically at gateway	P0
Monitor/Diagnosis/Planning/Safety agents	DOCUMENTED ONLY	Four bounded agents and outputs are specified ()	No proof outputs are computed from live state rather than templates	P1
Typed actions	DOCUMENTED ONLY	CREATE_REPLICA, SHIFT_ROUTE, MIGRATE_PARTITION, THROTTLE_WORKLOAD, ROLLBACK ()	Need schemas, validators, and rejection cases shown live	P0
Deterministic Policy Engine	DOCUMENTED ONLY	Allow/deny/modify/approval-required is required ()	No evidence decisions are deterministic or non-bypassable	P0
Decision Verifier	DOCUMENTED ONLY	Must validate state version, metrics, limits, permissions ()	No runtime proof verifier blocks unsafe or stale actions	P0
Action Gateway	DOCUMENTED ONLY	“Single execution path” and mandatory gateway ()	No evidence all mutation authority is centralized here	P0
Real runtime mutation	UNVERIFIED	PRD demands visible capacity/topology change and response in metrics ()	Could still be UI-only mutation without cluster effect	P0
Observability + audit lineage	DOCUMENTED ONLY	Prometheus/Grafana/Otel and audit lineage required ()	No proof logs are append-only or replayable	P1
Rollback/failure injection	DOCUMENTED ONLY	Injected execution failure should trigger rollback ()	No demonstrated rollback semantics or stored rollback material	P0

FIGURE 2 PRD compliance status for core ESA requirements

The remaining PRD items are weaker still. Rule-only baseline, ESA benchmark, intent interface, sensitive-data protection, and reproducibility are all clearly specified but still not evidenced as implemented in the supplied material (). The PRD is strong on what to prove, but it is not itself proof.

Anti-Mock Findings

The anti-mock posture must be strict because recent governance papers explicitly warn that architecture claims fail if execution still trusts the agent path or if authorization is never actually exercised (Fernandez, 2026; Rhodes & Kang, 2026).

Component	Likely status from supplied evidence	Why suspicious
Agent outputs	UNVERIFIED	No live traces or code show prompts, tool calls, or evidence-grounded reasoning
Diagnosis/planning	UNVERIFIED	Could be static JSON or templated action selection without live causal analysis
Policy results	UNVERIFIED	PRD requires allow, block, and stale reject, but no proof those cases occur ()
Runtime changes	UNVERIFIED	“Visible topology change” is required, not shown achieved ()
Rollback	UNVERIFIED	Rollback is a demo criterion, but no evidence of snapshot/compensation execution ()
Metrics	UNVERIFIED	Could be synthetic dashboard values disconnected from runtime
Audit records	DOCUMENTED ONLY	Schema exists, but append-only and replay semantics not proven ()
Stale-state rejection	DOCUMENTED ONLY	Central to thesis, but no test artifact proves it ()
Benchmarks	DOCUMENTED ONLY	Workload matrix and metrics are defined, but results absent ()

FIGURE 3 Anti-mock audit of unproven ESA components

Any direct shell, kubect!, or SQL path from agent code to runtime would be a serious architectural failure, because the attached PRD explicitly forbids unrestricted administrative execution by agents (). This matters because comparable industry and research systems already note that safe automation should mutate desired state through operators/controllers or controlled brokers, not by imperative agent commands (Bellala, 2025; Shevchenko, 2025; He & Yu, 2026).

Novelty And Razorpay

ESA duplicates established territory in five ways: MAPE-K control structure, multi-agent decomposition, intent-driven closed loops, observability-fed remediation, and payment-aware optimization contexts (Arcaini et al., 2015; Weyns et al., 2013; Gomes et al., 2021; Vadapalli, 2026; Bygari et al., 2021). Razorpay’s public smart-routing work already shows production AI improving payment success rates through feedback-updated routing decisions, so “AI plus payment infrastructure adaptation” alone is not new enough to attract curiosity (Bygari et al., 2021).

What ESA could uniquely add is a universal runtime abstraction where infrastructure proposals are typed, state-version-bound, deterministically admitted, executed only through one gateway, then measured and replayed without rerunning the model (Fatmi, 2026; SureshKumar, 2026; Rhodes & Kang, 2026). That is substantial if it is real, because much prior work still separates only policy from execution, or provides stateless evaluation, or stops at observability and remediation rather than provenance-complete governed execution (He & Yu, 2026; Fernandez, 2026; Rhodes & Kang, 2026).

Razorpay-specific relevance is therefore not “we built something Razorpay lacks.” Publicly, Razorpay already has production ML routing and a clear payment event surface through Test Mode and webhooks (Bygari et al., 2021). ESA remains interesting if it demonstrates how to keep probabilistic reasoning on the intent side while preserving deterministic authorization and auditable execution on the infrastructure side—the same governance tension highlighted in agentic payments literature (Davidovic & Tourpe, 2026; Amariles et al., 2026).

This is also where universality lives. The generic core is the closed-loop architecture, typed action contracts, policy/gateway boundary, version checks, measurement, rollback, and audit; the payment layer is just one adapter consisting of event types, payment workload semantics, webhook rules, and stricter correctness boundaries (Alqithami, 2026). Right now ESA reads closer to **D: multi-agent infrastructure controller** with payment framing than to **E: general adaptive runtime architecture**, because the public evidence does not yet show a domain-independent intermediate representation, reusable executor contracts, or portability demonstrated across non-payment substrates (Vadapalli, 2026; Lemos, 2023).

Safety, Gaps, And Plan

Research Extensions

Extension	Status	Value
Typed actions instead of arbitrary commands	DOCUMENTED ONLY	Very high
Deterministic authority	DOCUMENTED ONLY	Very high
State-version-aware action validity	DOCUMENTED ONLY	Very high
Safe no-op/rule fallback on model failure	DOCUMENTED ONLY	High
Human approval for high-risk actions	DOCUMENTED ONLY	High
Expected effect vs observed effect	PARTIAL in PRD concept	High
Decision replay without original LLM call	DOCUMENTED ONLY	High
Intent → measurable constraints	DOCUMENTED ONLY	Medium-high
Closed-loop adaptation after execution	PARTIAL in PRD concept	Medium-high
Executable state = data + policy + constraints + metadata	PARTIAL conceptually	Medium-high

FIGURE 4 Status of highest-value ESA research extensions

The core safety question is simple: can the agent bypass the deterministic path? The attached PRD says no agent may execute raw admin commands, all execution must go through one gateway, stale versions must be rejected, high-risk actions may require approval, and model failure should fall back to rules (). But without code or runtime evidence, those are **claims of intended authority separation**, not verified invariants.

The benchmark standard also has to be higher than attractive screenshots. Agent-evaluation literature repeatedly finds poor reproducibility, weak execution semantics, and protocol incomparability in closed-loop agent systems, while self-healing surveys note reproducibility across heterogeneous environments remains open (Xia et al., 2026; Ganapathi & Gupta, 2026). ESA’s own benchmark design is sound on paper—same workloads, rule-only baseline, before/after metrics, ablations—but it is not yet scientifically credible until multiple repeatable runs and raw traces exist ().

Top Changes

Rank	Change	Why
1	Prove Action Gateway is the only mutation path	Makes authority claim real
2	Demo one stale-state rejection live	Best differentiator against standard AIOps
3	Show real topology mutation plus metric response	Kills UI-only skepticism
4	Add append-only replayable audit artifact	Supports defensibility and postmortem trust
5	Inject one execution failure and successful rollback	Proves safe failure
6	Show one policy allow and one deny	Exercises governance boundary
7	Expose typed action schemas and validator logs	Distinguishes from raw tool use
8	Add rule-only vs ESA repeated benchmark runs	Prevents cherry-pick skepticism
9	Separate payment adapters from generic core in code	Strengthens universality claim
10	Prove model timeout falls back to no-op/rule path	Critical safety assurance

FIGURE 5 Prioritized technical changes for ESA credibility

CRITICAL GAPS: no runtime proof of non-bypassable authority, stale-state rejection, rollback, or real mutation. **HIGH-VALUE GAPS:** replayable decision records, expected-vs-observed effect checks, generic action IR beyond payment scenarios. **PRESENTATION GAPS:** hidden evidence path from event to metric recovery, no explicit deny case, no visibly bounded agent permissions. **UNIVERSALITY/RESEARCH GAPS:** “executable state” remains mostly a concept until state drives verification and execution semantics across multiple domains (Qin et al., 2026; Kosińska et al., 2026).

The most important shift in the evidence is that the hard problem is no longer building agentic control loops, but building a deterministic execution boundary around them (Sohail & Haider, 2026; Alqithami, 2026; Fatmi, 2026). The biggest open question is whether ESA’s current repository actually enforces that boundary in code and runtime, or only describes it convincingly in the attached PRD.

THE 5 THINGS WE MUST FIX

1. Make the Action Gateway the only mutation path and prove it in logs.
2. Demonstrate one live stale-state rejection at the gateway.
3. Show real replica or route mutation with before/after metrics.
4. Add append-only replayable audit artifacts with state and policy versions.
5. Inject one execution failure and prove rollback restores a valid prior state.

THE 5 THINGS WE SHOULD NOT BUILD

1. More agents.
2. More UI polish.

3. Payment-business features beyond the hotspot demo.
4. Predictive scaling research before safety proof.
5. New subsystems after the core loop still lacks runtime evidence.

THE ONE THING THAT MAKES ESA INTERESTING

A model can propose infrastructure actions, but only a deterministic, state-version-aware, typed-action runtime is allowed to turn them into real mutations.

THE ONE SENTENCE RAZORPAY SHOULD REMEMBER

ESA is only interesting if it proves that agent reasoning can shape infrastructure decisions without ever holding infrastructure authority.

UNIVERSAL ESA FORMULATION

INTENT → CONSTRAINTS → EXECUTABLE STATE → OBSERVE → REASON → TYPED ACTION → POLICY → VERIFY → EXECUTE → MEASURE → UPDATE → AUDIT/ROLLBACK

The generic runtime architecture is intent compilation, state representation, observation, reasoning, typed action generation, deterministic policy and verification, controlled execution, measurement, update, and audit/rollback (SureshKumar, 2026). The domain adapters are payment-event ingestion, webhook validation, pseudonymized payment metadata, payment correctness constraints, and payment-specific risk/rollback semantics ().

These search results were found and analyzed using Consensus, an AI-powered search engine for research. Try it at <https://consensus.app>. © 2026 Consensus NLP, Inc. Personal, non-commercial use only; redistribution requires copyright holders' consent.

References

ESA | Payment-Angle Product Requirements Document Executable State Architecture (ESA) | Detailed Payment Infrastructure Prototype PRD.

Alqithami, S. (2026). Autonomous Agents on Blockchains: Standards, Execution Models, and Trust Boundaries. *ArXiv*, abs/2601.04583. <https://doi.org/10.48550/arxiv.2601.04583>

Amariles, D. R., Charlotin, D., & He-Guelton, L. (2026). AI Agents in Payments: Applications, Risks and Regulations. *European Journal of Risk Regulation*. <https://doi.org/10.1017/err.2026.10103>

Aravindakshan, V. (2026). An Agentic AI Framework for Observability-Driven Self-Healing and Auto-Remediation in Cloud Infrastructure. *2026 9th International Conference on Inventive Computation Technologies (ICICT)*, 1781-1786. <https://doi.org/10.1109/iciict68280.2026.11510957>

Arcaini, P., Riccobene, E., & Scandurra, P. (2015). Modeling and Analyzing MAPE-K Feedback Loops for Self-Adaptation. *2015 IEEE/ACM 10th International Symposium on Software Engineering for Adaptive and Self-Managing Systems*, 13-23. <https://doi.org/10.1109/seams.2015.10>

Avgerinos, V., Ramantas, K., Alonso, L., & Verikoukis, C. (2026). ARM: Autonomous Remediation and Management With LLM Agents for Intent-Driven Control. *IEEE Internet of Things Journal*, 13, 18305-18315. <https://doi.org/10.1109/jiot.2025.3648858>

Bellala, K. R. (2025). Towards Autonomous Kubernetes: A Framework for AI-Driven Operations (AIOps). *International Journal of Innovative Science and Research Technology*. <https://doi.org/10.38124/ijisrt/25sep1016>

Bygari, R., Gupta, A., Raghuvanshi, S., Bapna, A., & Sahu, B. (2021). An AI-powered Smart Routing Solution for Payment Systems. *2021 IEEE International Conference on Big Data (Big Data)*, 2026-2033. <https://doi.org/10.1109/bigdata52589.2021.9671961>

Davidovic, S., & Tourpe, H. (2026). How Agentic AI Will Reshape Payments. *IMF Notes*. <https://doi.org/10.5089/9781513533308.068>

- Fatmi, A. (2026). Faramesh: A Protocol-Agnostic Execution Control Plane for Autonomous Agent Systems. *ArXiv*, *abs/2601.17744*. <https://doi.org/10.48550/arxiv.2601.17744>
- Fernandez, M. (2026). Agent Control Protocol: Admission Control for Agent Actions. *ArXiv*, *abs/2603.18829*. <https://doi.org/10.5281/zenodo.19672575>
- Ganapathi, A., & Gupta, M. (2026). Autonomous AI Agents for Self-Healing Cloud Infrastructure: A Survey with Critical Evaluation. *2026 International Conference on Integrated Intelligence and Cognitive Engineering (ICIICE)*, 1-6. <https://doi.org/10.1109/iciice69672.2026.11565093>
- Gomes, P., Buhrgard, M., Harmatos, J., Mohalik, S., Roeland, D., & Niemöller, J. (2021). Intent-driven Closed Loops for Autonomous Networks. *J. ICT Stand.*, 9, 257-290. <https://doi.org/10.13052/jicts2245-800x.929>
- He, J., & Yu, D. (2026). Sovereign Execution Broker: Enforcing Certificate-Bound Authority in Agentic Control Planes. *ArXiv*, *abs/2606.20520*. <https://doi.org/10.48550/arxiv.2606.20520>
- Kosińska, J., Saków, K., & Szczurek, M. (2026). An Execution Controller Enabling Autonomic Features in the Cloud-Native Environment. *ACM Transactions on Internet Technology*, 26, 1 - 24. <https://doi.org/10.1145/3799702>
- Lemos, R. D. (2023). Software Self-adaptation and Industry: Blame MAPE-K. *2023 IEEE/ACM 18th Symposium on Software Engineering for Adaptive and Self-Managing Systems (SEAMS)*, 88-89. <https://doi.org/10.1109/seams59076.2023.00021>
- Qin, X., Luan, S., See, J., Yang, C., & Li, Z. (2026). ECM Contracts: Contract-Aware, Versioned, and Governable Capability Interfaces for Embodied Agents. *ArXiv*, *abs/2604.13097*. <https://doi.org/10.48550/arxiv.2604.13097>
- Rhodes, J., & Kang, G. (2026). Proof of Execution: Runtime Verification for Governed AI Agent Actions. <https://doi.org/10.48550/arxiv.2607.05397>
- Sannareddy, S. B. (2025). Autonomous Kubernetes Cluster Healing using Machine Learnin. *International Journal of Research Publications in Engineering, Technology and Management*. <https://doi.org/10.15662/ijrpetm.2024.0705006>
- Shevchenko, O. (2025). Towards Self-Healing Cloud Infrastructure: Automated Recovery Methods and Their Effectiveness. *The American Journal of Engineering and Technology*. <https://doi.org/10.37547/tajet/volume07issue06-10>
- Sohail, S., & Haider, G. (2026). Bounded Autonomy for Enterprise AI: Typed Action Contracts and Consumer-Side Execution. *ArXiv*, *abs/2604.14723*. <https://doi.org/10.48550/arxiv.2604.14723>
- SureshKumar, S. (2026). R-LAM: Reproducibility-Constrained Large Action Models for Scientific Workflow Automation. *ArXiv*, *abs/2601.09749*. <https://doi.org/10.48550/arxiv.2601.09749>
- Vadapalli, V. L. N. K. (2026). Agentic Infrastructure as Code: A Closed-Loop Framework for Autonomous and Self-Optimizing Cloud Systems. *International Journal For Multidisciplinary Research*. <https://doi.org/10.36948/ijfmr.2026.v08i02.78058>
- Weyns, D., Schmerl, B., Grassi, V., Malek, S., Mirandola, R., Prehofer, C., Wuttke, J., Andersson, J., Giese, H., & Göschka, K. M. (2013). On Patterns for Decentralized Control in Self-Adaptive Systems. 76-107. https://doi.org/10.1007/978-3-642-35813-5_4
- Xia, Y., You, P., Wang, T., Liu, F., Qi, H., Wu, S. X., & Zhang, S. (2026). Agentic Trading: When LLM Agents Meet Financial Markets. *ArXiv*, *abs/2605.19337*. <https://doi.org/10.48550/arxiv.2605.19337>